

December 3, 2024

1 Exploration data analysis

This final project for CSCA5622 Introduction to machine learning: Supervised Learning. We are using data from UCI Machine Learning Repository ([UCI Repository](#)). In this opportunity, we are using the dataset called ‘Bank Marketing’, the data is related with direct marketing campaigns (phone calls) of a Portuguese banking institution. The classification goal is to predict if the client will subscribe a term deposit (variable y).

2 Dataset information

The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be (‘yes’) or not (‘no’) subscribed.

2.1 Variables tables

Variable Name	Role	Type	Demographic	Description	Units	Missing Values
age	Feature	Integer	Age	Age of the client		no
job	Feature	Categorical	Occupation	Type of job		no
marital	Feature	Categorical	Marital Status	Marital status		no
education	Feature	Categorical	Education Level	Education level		no
default	Feature	Binary		Whether the client has credit in default		no
balance	Feature	Integer		Average yearly balance in euros		no
housing	Feature	Binary		Whether the client has a housing loan		no
loan	Feature	Binary		Whether the client has a personal loan		no
contact	Feature	Categorical		Communication type used to contact the client		yes
day_of_week	Feature	Date		Last contact day of the week		no
month	Feature	Date		Last contact month of the year		no
duration	Feature	Integer		Duration of the last contact in seconds (numeric).		no
campaign	Feature	Integer		Number of contacts performed during this campaign for this client		no

Variable Name	Role	Type	Demographic	Description	Units	Missing Values
pdays	Feature	Integer		Number of days since the client was last contacted from a previous campaign		yes
previous	Feature	Integer		Number of contacts performed before this campaign and for this client		no
poutcome	Feature	Categorical		Outcome of the previous marketing campaign ('failure', 'nonexistent', 'success')		yes
y	Target	Binary		Whether the client subscribed to a term deposit		no

2.2 Additional information

2.2.1 Input variables:

Bank client data: - **age** (numeric) - **job** : type of job (categorical: "admin.", "unknown", "unemployed", "management", "housemaid", "entrepreneur", "student", "blue-collar", "self-employed", "retired", "technician", "services") - **marital** : marital status (categorical: "married", "divorced", "single"; note: "divorced" means divorced or widowed) - **education** (categorical: "unknown", "secondary", "primary", "tertiary") - **default**: has credit in default? (binary: "yes", "no") - **balance**: average yearly balance, in euros (numeric) - **housing**: has housing loan? (binary: "yes", "no") - **loan**: has personal loan? (binary: "yes", "no") #### Related with the last contact of the current campaign: - **contact**: contact communication type (categorical: "unknown", "telephone", "cellular") - **day**: last contact day of the month (numeric) - **month**: last contact month of year (categorical: "jan", "feb", "mar", ..., "nov", "dec") - **duration**: last contact duration, in seconds (numeric) #### Other attributes: - **campaign**: number of contacts performed during this campaign and for this client (numeric, includes last contact) - **pdays**: number of days that passed by after the client was last contacted from a previous campaign (numeric, -1 means client was not previously contacted) - **previous**: number of contacts performed before this campaign and for this client (numeric) - **poutcome**: outcome of the previous marketing campaign (categorical: "unknown", "other", "failure", "success")

2.2.2 Output variable (desired target):

- **y** - has the client subscribed a term deposit? (binary: "yes", "no")

3 Project Summary

3.1 Objective

The goal of this project is to build a classification model that predicts whether a customer will subscribe to a term deposit based on various features from the bank's marketing campaign data. The project will involve exploring and preprocessing the data, selecting relevant features, applying machine learning algorithms, and evaluating model performance using appropriate metrics. The final aim is to develop a robust model that helps the bank optimize its marketing efforts by targeting customers who are most likely to subscribe to a term deposit.

3.2 Import libraries

```
[24]: from ucimlrepo import fetch_ucirepo
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from utils import *

warnings.filterwarnings('ignore')
pd.set_option('display.max_columns',None)
pd.set_option('display.max_rows',None)

# fetch dataset
bank_marketing = fetch_ucirepo(id=222)

# data (as pandas dataframes)
X = bank_marketing.data.features
y = bank_marketing.data.targets

# metadata
print(bank_marketing.metadata)

# variable information
print(bank_marketing.variables)
```

```
{'uci_id': 222, 'name': 'Bank Marketing', 'repository_url':
'https://archive.ics.uci.edu/dataset/222/bank+marketing', 'data_url':
'https://archive.ics.uci.edu/static/public/222/data.csv', 'abstract': 'The data
is related with direct marketing campaigns (phone calls) of a Portuguese banking
institution. The classification goal is to predict if the client will subscribe
a term deposit (variable y).', 'area': 'Business', 'tasks': ['Classification'],
'characteristics': ['Multivariate'], 'num_instances': 45211, 'num_features': 16,
'feature_types': ['Categorical', 'Integer'], 'demographics': ['Age',
'Occupation', 'Marital Status', 'Education Level'], 'target_col': ['y'],
'index_col': None, 'has_missing_values': 'yes', 'missing_values_symbol': 'NaN',
'year_of_dataset_creation': 2014, 'last_updated': 'Fri Aug 18 2023',
'dataset_doi': '10.24432/C5K306', 'creators': ['S. Moro', 'P. Rita', 'P.
Cortez'], 'intro_paper': {'ID': 277, 'type': 'NATIVE', 'title': 'A data-driven
approach to predict the success of bank telemarketing', 'authors': 'Sérgio Moro,
P. Cortez, P. Rita', 'venue': 'Decision Support Systems', 'year': 2014,
'journal': None, 'DOI': '10.1016/j.dss.2014.03.001', 'URL': 'https://www.semanti
cscholar.org/paper/cab86052882d126d43f72108c6cb41b295cc8a9e', 'sha': None,
'corpus': None, 'arxiv': None, 'mag': None, 'acl': None, 'pmid': None, 'pmcid':
None}, 'additional_info': {'summary': "The data is related with direct marketing
campaigns of a Portuguese banking institution. The marketing campaigns were
```

based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed. \n\nThere are four datasets: \n1) bank-additional-full.csv with all examples (41188) and 20 inputs, ordered by date (from May 2008 to November 2010), very close to the data analyzed in [Moro et al., 2014]\n2) bank-additional.csv with 10% of the examples (4119), randomly selected from 1), and 20 inputs.\n3) bank-full.csv with all examples and 17 inputs, ordered by date (older version of this dataset with less inputs). \n4) bank.csv with 10% of the examples and 17 inputs, randomly selected from 3 (older version of this dataset with less inputs). \nThe smallest datasets are provided to test more computationally demanding machine learning algorithms (e.g., SVM). \n\nThe classification goal is to predict if the client will subscribe (yes/no) a term deposit (variable y).\", 'purpose': None, 'funded_by': None, 'instances_represent': None, 'recommended_data_splits': None, 'sensitive_data': None, 'preprocessing_description': None, 'variable_info': 'Input variables:\n# bank client data:\n 1 - age (numeric)\n 2 - job : type of job (categorical: \"admin.\",\"unknown\",\"unemployed\",\"management\",\"housemaid\",\"entrepreneur\",\"student\", \"blue-collar\",\"self-employed\",\"retired\",\"technician\",\"services\") \n 3 - marital : marital status (categorical: \"married\",\"divorced\",\"single\"; note: \"divorced\" means divorced or widowed)\n 4 - education (categorical: \"unknown\",\"secondary\",\"primary\",\"tertiary\")\n 5 - default: has credit in default? (binary: \"yes\",\"no\")\n 6 - balance: average yearly balance, in euros (numeric) \n 7 - housing: has housing loan? (binary: \"yes\",\"no\")\n 8 - loan: has personal loan? (binary: \"yes\",\"no\")\n # related with the last contact of the current campaign:\n 9 - contact: contact communication type (categorical: \"unknown\",\"telephone\",\"cellular\") \n 10 - day: last contact day of the month (numeric)\n 11 - month: last contact month of year (categorical: \"jan\", \"feb\", \"mar\", ..., \"nov\", \"dec\")\n 12 - duration: last contact duration, in seconds (numeric)\n # other attributes:\n 13 - campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)\n 14 - pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric, -1 means client was not previously contacted)\n 15 - previous: number of contacts performed before this campaign and for this client (numeric)\n 16 - poutcome: outcome of the previous marketing campaign (categorical: \"unknown\",\"other\",\"failure\",\"success\")\n\nOutput variable (desired target):\n 17 - y - has the client subscribed a term deposit? (binary: \"yes\",\"no\")\n', 'citation': None}}

	name	role	type	demographic \
0	age	Feature	Integer	Age
1	job	Feature	Categorical	Occupation
2	marital	Feature	Categorical	Marital Status
3	education	Feature	Categorical	Education Level
4	default	Feature	Binary	None
5	balance	Feature	Integer	None
6	housing	Feature	Binary	None
7	loan	Feature	Binary	None
8	contact	Feature	Categorical	None

9	day_of_week	Feature	Date	None
10	month	Feature	Date	None
11	duration	Feature	Integer	None
12	campaign	Feature	Integer	None
13	pdays	Feature	Integer	None
14	previous	Feature	Integer	None
15	poutcome	Feature	Categorical	None
16	y	Target	Binary	None

		description	units	missing_values	
0			None	None	no
1	type of job (categorical: 'admin.', 'blue-colla...		None		no
2	marital status (categorical: 'divorced', 'marri...		None		no
3	(categorical: 'basic.4y', 'basic.6y', 'basic.9y'...		None		no
4		has credit in default?	None		no
5		average yearly balance	euros		no
6		has housing loan?	None		no
7		has personal loan?	None		no
8	contact communication type (categorical: 'cell...		None		yes
9		last contact day of the week	None		no
10	last contact month of year (categorical: 'jan'...		None		no
11	last contact duration, in seconds (numeric). ...		None		no
12	number of contacts performed during this campa...		None		no
13	number of days that passed by after the client...		None		yes
14	number of contacts performed before this campa...		None		no
15	outcome of the previous marketing campaign (ca...		None		yes
16		has the client subscribed a term deposit?	None		no

```
[25]: # Join the feature dataset and the target dataset to start the exploratory data
      ↪analysis
df = X.copy()
df['y'] = y
df.sample(3)
```

```
[25]:      age      job  marital education default  balance housing loan \
38259   45 self-employed  married  tertiary     no    -497     yes  yes
15526   35 entrepreneur  single  tertiary     no     145     yes   no
42082   61      retired  married  primary     no    8729     no   no

      contact  day_of_week month  duration  campaign  pdays  previous \
38259  cellular         15   may      217         1    176         1
15526  cellular         18   jul      799         2     -1         0
42082  cellular         30  oct      480         1     -1         0

      poutcome  y
38259  failure  no
15526     NaN   yes
```

42082 NaN yes

4 Validation of null values and duplicates rows

```
[26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             45211 non-null  int64
1   job             44923 non-null  object
2   marital         45211 non-null  object
3   education       43354 non-null  object
4   default         45211 non-null  object
5   balance         45211 non-null  int64
6   housing         45211 non-null  object
7   loan            45211 non-null  object
8   contact         32191 non-null  object
9   day_of_week     45211 non-null  int64
10  month           45211 non-null  object
11  duration        45211 non-null  int64
12  campaign        45211 non-null  int64
13  pdays           45211 non-null  int64
14  previous        45211 non-null  int64
15  poutcome       8252 non-null   object
16  y               45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

```
[27]: #Percentage of null values per column
```

```
print(df.isna().sum()/df.shape[0]*100)
```

```
age             0.000000
job             0.637013
marital         0.000000
education       4.107407
default         0.000000
balance         0.000000
housing         0.000000
loan            0.000000
contact         28.798301
day_of_week     0.000000
month           0.000000
duration        0.000000
```

```

campaign      0.000000
pdays        0.000000
previous      0.000000
poutcome      81.747805
y             0.000000
dtype: float64

```

We observe that in the dataset, the variables with the highest number of null values are the following:

- poutcome : 81.74%
- contact : 28.79
- education : 4.10%
- job : 0.63%

Therefore, we must define the strategy according to the context of each variable. For the categorical variables, poutcome, contact, education we will fill the empty fields with the category “unknown”, since the variable definition allows us to use that category when the data is unknown. For the job variable, we can have two strategies, either we delete the rows with empty fields, or we fill the fields with the fashion of the category, since the percentage of null values is very small, we will perform the second strategy.

```

[28]: df['poutcome'] = df['poutcome'].fillna('unknown')
      df['contact'] = df['contact'].fillna('unknown')
      df['education'] = df['education'].fillna('unknown')
      df['job'] = df['job'].fillna(df['job'].mode().values[0])

```

```

[29]: print(df.isna().sum()/df.shape[0]*100)

```

```

age          0.0
job          0.0
marital      0.0
education    0.0
default      0.0
balance      0.0
housing      0.0
loan         0.0
contact      0.0
day_of_week  0.0
month        0.0
duration     0.0
campaign     0.0
pdays       0.0
previous     0.0
poutcome     0.0
y            0.0
dtype: float64

```

```

[30]: # Selecting duplicate rows except first
      # occurrence based on all columns

```

```

duplicate = df[df.duplicated()]

print("Duplicate Rows :")

# Print the resultant Dataframe
duplicate

```

Duplicate Rows :

```

[30]: Empty DataFrame
Columns: [age, job, marital, education, default, balance, housing, loan,
contact, day_of_week, month, duration, campaign, pdays, previous, poutcome, y]
Index: []

```

We don't have duplicate values in the dataset. So we don't need to drop rows in this case.

5 Statistical Analysis

```

[31]: df.describe().T

```

```

[31]:
count      mean      std      min      25%      50%      75%  \
age      45211.0    40.936210   10.618762   18.0    33.0    39.0    48.0
balance   45211.0  1362.272058  3044.765829 -8019.0    72.0   448.0   1428.0
day_of_week 45211.0   15.806419    8.322476    1.0     8.0   16.0    21.0
duration   45211.0   258.163080  257.527812    0.0   103.0  180.0   319.0
campaign   45211.0    2.763841    3.098021    1.0     1.0    2.0     3.0
pdays     45211.0   40.197828  100.128746   -1.0    -1.0   -1.0    -1.0
previous   45211.0    0.580323    2.303441    0.0     0.0    0.0     0.0

max
age      95.0
balance  102127.0
day_of_week 31.0
duration  4918.0
campaign  63.0
pdays   871.0
previous 275.0

```

6 Univariable analysis

6.1 Categorical variables

```

[32]: # Categorical variables
categorical_columns = [
    'job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'poutcome']

# Create subplots for each variable

```



```

fig, axes = plt.subplots(nrows=2, ncols=5, figsize=(30, 15), sharey=False)

# Flatten the axes array for easier indexing
axes = axes.flatten()

# Plot each categorical column
for i, column in enumerate(categorical_columns):
    percentage = df[column].value_counts(normalize=True) * 100
    percentage_df = percentage.reset_index()
    percentage_df.columns = [column, 'Percentage']

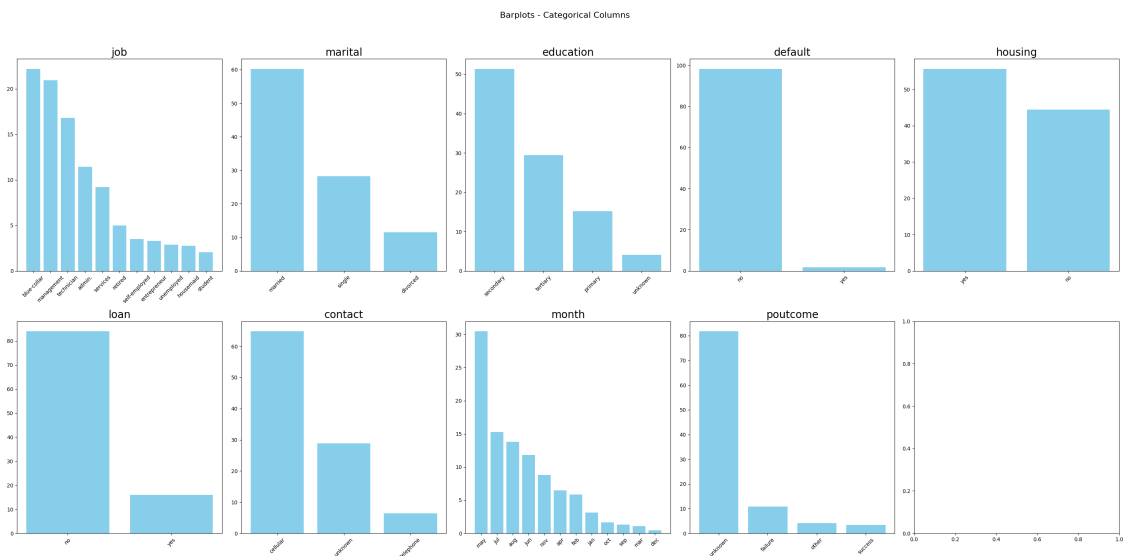
    axes[i].bar(percentages_df[column], percentages_df['Percentage'],
        color='skyblue')
    axes[i].set_title(column, fontsize=20)
    axes[i].tick_params(axis='x', rotation=45) # Rotate x-axis labels for
        better visibility

# Add a main title
fig.suptitle("Barplots - Categorical Columns", fontsize=16)

# Adjust layout to prevent overlap
plt.tight_layout(rect=[0, 0, 1, 0.95])

# Show the plot
plt.show()

```



```
[33]: df['y'].value_counts(normalize=True)*100
```

```
[33]: y
      no      88.30152
      yes     11.69848
      Name: proportion, dtype: float64
```

```
[34]: df['default'].value_counts(normalize=True)*100
```

```
[34]: default
      no      98.197341
      yes      1.802659
      Name: proportion, dtype: float64
```

On this occasion, we observed that the “default” variable has very little variability (98% no, 1.8% yes). Therefore, we see that the information provided by the categorical variable is not very significant, and we could even consider deleting the variable when training the model. Therefore, we are going to delete this variable from the final dataset.

6.2 Numerical variables

6.3 Outliers

```
[35]: numerical_columns = df.select_dtypes([np.number]).columns.tolist()
      print(numerical_columns)
      numeric_statistics(data=df,numeric_columns=numerical_columns)
```

```
['age', 'balance', 'day_of_week', 'duration', 'campaign', 'pdays', 'previous']
```

```
[35]:
```

	Variable	Median	Mean	Std. Deviation	Min. Value	Percentile 5 \
0	age	39.0	40.936210	10.618762	18	27.0
1	balance	448.0	1362.272058	3044.765829	-8019	-172.0
2	day_of_week	16.0	15.806419	8.322476	1	3.0
3	duration	180.0	258.163080	257.527812	0	35.0
4	campaign	2.0	2.763841	3.098021	1	1.0
5	pdays	-1.0	40.197828	100.128746	-1	-1.0
6	previous	0.0	0.580323	2.303441	0	0.0

	Percentile 10	Percentile 15	Percentile 20	Percentile 25	Percentile 50 \
0	29.0	30.0	32.0	33.0	39.0
1	0.0	0.0	22.0	72.0	448.0
2	5.0	6.0	7.0	8.0	16.0
3	58.0	75.0	89.0	103.0	180.0
4	1.0	1.0	1.0	1.0	2.0
5	-1.0	-1.0	-1.0	-1.0	-1.0
6	0.0	0.0	0.0	0.0	0.0

	Percentile 75	Percentile 80	Percentile 85	Percentile 90	Percentile 95 \
0	48.0	51.0	53.0	56.0	59.0
1	1428.0	1859.0	2539.0	3574.0	5768.0
2	21.0	24.0	27.0	28.0	29.0

3	319.0	368.0	437.0	548.0	751.0
4	3.0	4.0	4.0	5.0	8.0
5	-1.0	-1.0	102.0	185.0	317.0
6	0.0	0.0	1.0	2.0	3.0


	Percentile 99	Max. value
0	71.0	95
1	13164.9	102127
2	31.0	31
3	1269.0	4918
4	16.0	63
5	370.0	871
6	8.9	275

```
[36]: # Create subplots for each variable on its own scale
fig, axes = plt.subplots(nrows=1, ncols=7, figsize=(20, 6), sharey=False)

# Plot each column separately
for i, column in enumerate(numerical_columns):
    axes[i].boxplot(df[column], vert=True)
    axes[i].set_title(column, fontsize=10)
    axes[i].tick_params(axis='x', which='both', bottom=False, top=False,
        labelbottom=False)
    axes[i].tick_params(axis='y', labelsize=8)

# Add a main title
fig.suptitle("Boxplots - Numerical columns", fontsize=16)
plt.tight_layout(rect=[0, 0, 1, 0.95])

# Show the plot
plt.show()
```



We can observe that in the variables age, balance, duration, campaign, pdays, previous we have outliers to identify. If we do not refine these values, we will not be able to have an optimal model that generalizes the results when tested with other populations or samples. On this occasion, we

are going to use the winzor approach.

```
[37]: from scipy.stats.mstats import winsorize


outlier_columns = ['age', 'balance', 'duration', 'campaign', 'pdays', 'previous']
for col in outlier_columns:
    data = df[col].values
    winzorized_data = winsorize(data, limits=[0.05, 0.05])
    df[col] = winzorized_data
```

```
[38]: # Create subplots for each variable on its own scale
fig, axes = plt.subplots(nrows=1, ncols=7, figsize=(20, 6), sharey=False)

# Plot each column separately
for i, column in enumerate(numerical_columns):
    axes[i].boxplot(df[column], vert=True)
    axes[i].set_title(column, fontsize=10)
    axes[i].tick_params(axis='x', which='both', bottom=False, top=False,
        ↪labelbottom=False)
    axes[i].tick_params(axis='y', labelsize=8)

# Add a main title
fig.suptitle("Boxplots - Numerical columns", fontsize=16)
plt.tight_layout(rect=[0, 0, 1, 0.95])

# Show the plot
plt.show()
```




7 Correlation Analysis

In this part, we start with the correlation analysis between the categorical and numerical variables with the target variable ('y'). To determine the correlation and interaction between the variables of the model and the target variable, with this we will be able to define which variables we will finally take in the model, since we will need to consider the variables that have more information and interaction with the target variable. Subsequently, we will review the interaction between

numerical and categorical variables, to validate that there are no variables that give us a high interaction and may cause the model to take redundant variables.

```
[39]: relationship_with_target(df=df)
```

Analyzing age vs Target (y):




y=no: Mean=40.75, Std=9.51

y=yes: Mean=40.88, Std=10.99


Analyzing job vs Target (y):

Chi-square Test between job and y: p-value = 5.575427995540736e-172




Analyzing marital vs Target (y):

Chi-square Test between marital and y: p-value = 2.1450999986791792e-43




Analyzing education vs Target (y):

Chi-square Test between education and y: p-value = $1.6266562124072994e-51$

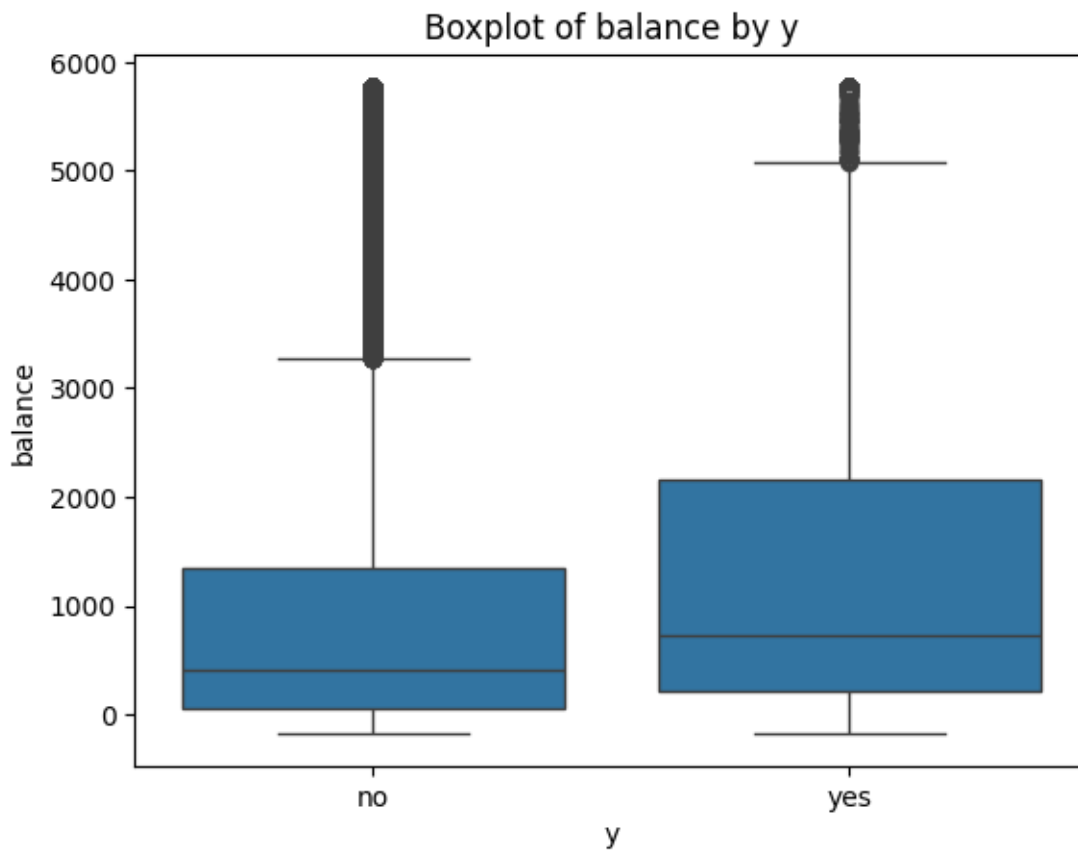


Analyzing default vs Target (y):

Chi-square Test between default and y: p-value = 2.4538606753508344e-06



Analyzing balance vs Target (y):




y=no: Mean=1070.83, Std=1556.99

y=yes: Mean=1465.41, Std=1714.48


Analyzing housing vs Target (y):

Chi-square Test between housing and y: p-value = 2.918797605076633e-192




Analyzing loan vs Target (y):

Chi-square Test between loan and y: p-value = 1.665061163492756e-47

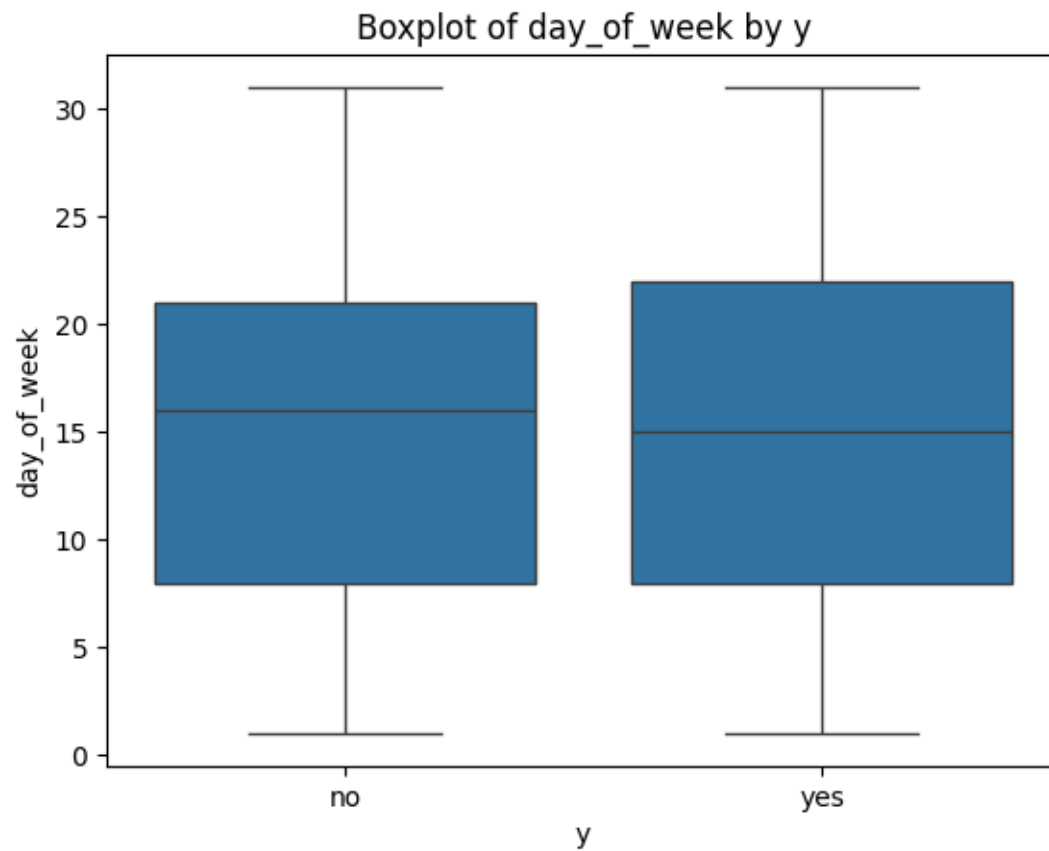


Analyzing contact vs Target (y):

Chi-square Test between contact and y: p-value = 1.251738325340638e-225




Analyzing day_of_week vs Target (y):




y=no: Mean=15.89, Std=8.29
y=yes: Mean=15.16, Std=8.50

Analyzing month vs Target (y):
Chi-square Test between month and y: p-value = 0.0




Analyzing duration vs Target (y):



y=no: Mean=214.18, Std=167.40

y=yes: Mean=453.89, Std=228.83

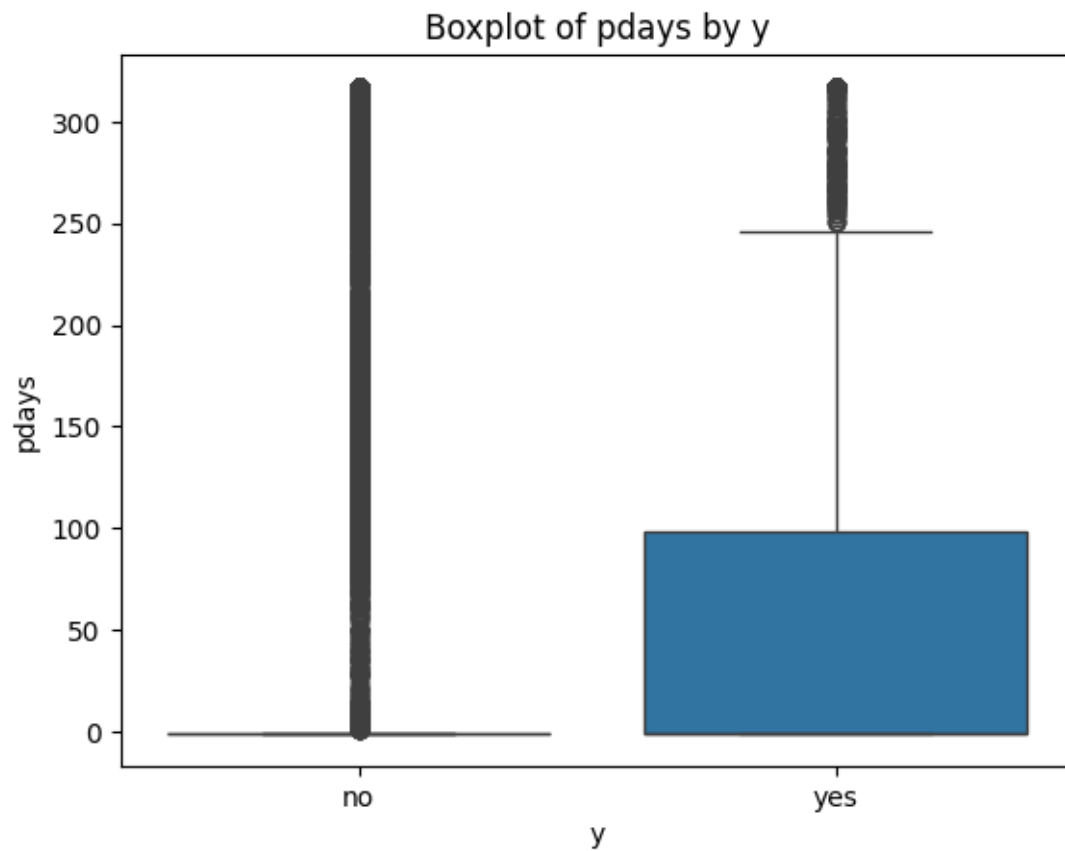
Analyzing campaign vs Target (y):



y=no: Mean=2.57, Std=1.94

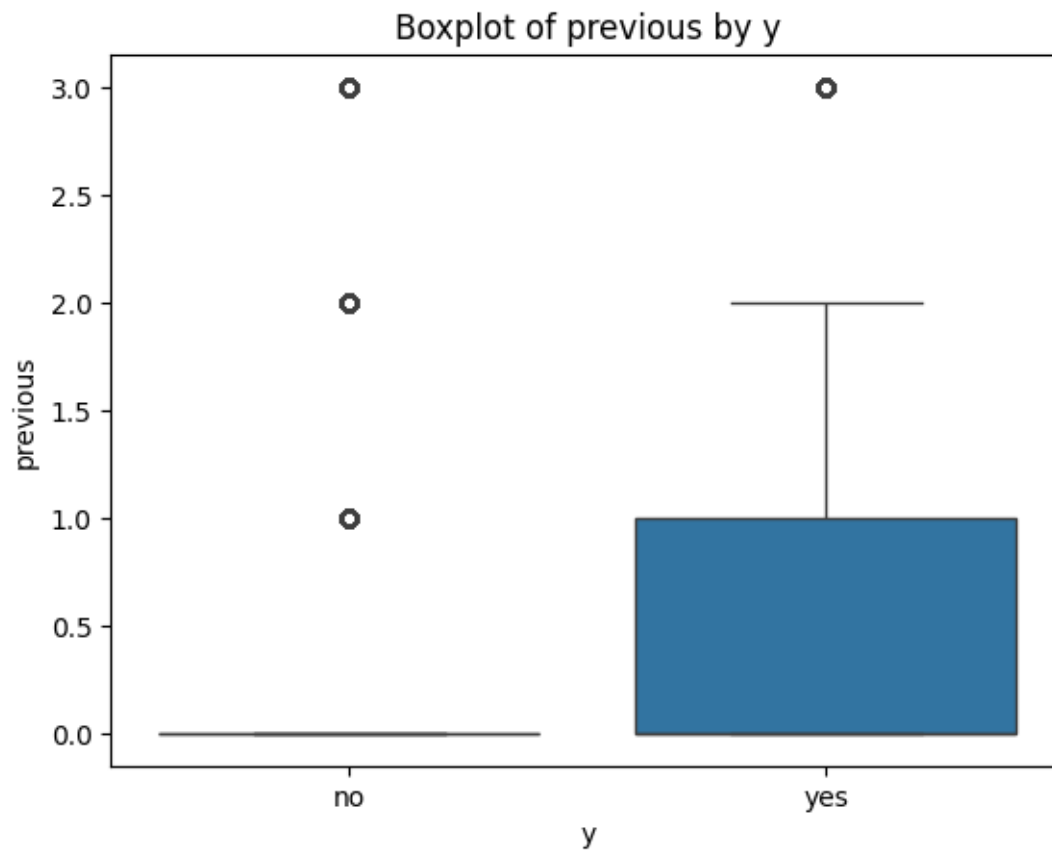
y=yes: Mean=2.07, Std=1.53

Analyzing pdays vs Target (y):




y=no: Mean=34.20, Std=88.50
y=yes: Mean=62.89, Std=99.87

Analyzing previous vs Target (y):




y=no: Mean=0.33, Std=0.82
y=yes: Mean=0.77, Std=1.15

Analyzing poutcome vs Target (y):
Chi-square Test between poutcome and y: p-value = 0.0



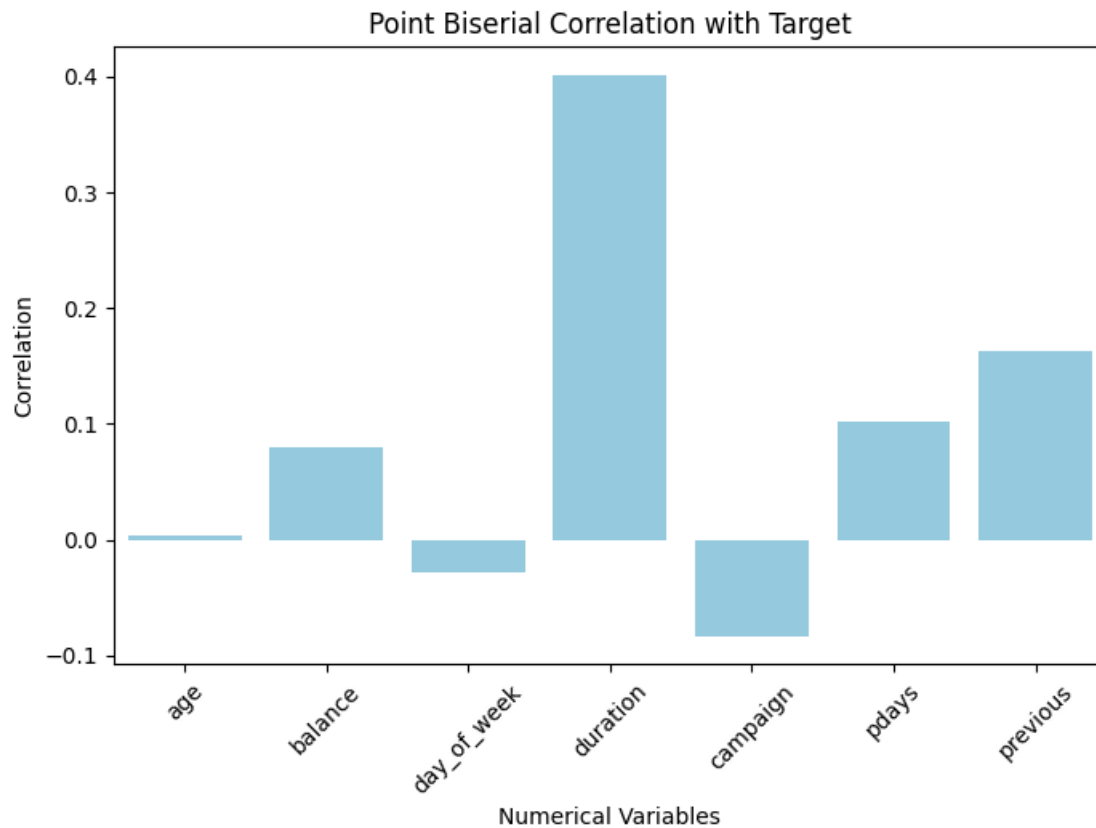
With the observed results, we conclude that the “default” variable does not influence the “y” objective, so it may be detrimental to add a variable that does not contribute relevant information to the model. For the numerical variables we observe that in all variables there is a difference in the distribution between the categories of the target variable “y”. Therefore, we will keep all the numerical variables for the final dataset of the model.

```
[40]: # Compute Cramér's V matrix
cramers_matrix = cramers_v_matrix(df, categorical_columns + ['y'])
plt.figure(figsize=(10, 8))
sns.heatmap(cramers_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Cramér's V Matrix")
plt.show()
```




In the cramer's matrix, we note that the categorical variables do not have a strong interaction between them, however we observe that the categorical variable poutcome, month are the variables with the highest correlation with the target 'y'.


```
[41]: # Compute the matrix
biserial_matrix = point_biserial_matrix(df, numerical_columns, 'y')
plt.figure(figsize=(8, 5))
sns.barplot(x=biserial_matrix.index, y="Point Biserial Correlation",
            data=biserial_matrix.reset_index(), color='skyblue')
plt.xticks(rotation=45)
plt.title("Point Biserial Correlation with Target")
plt.ylabel("Correlation")
plt.xlabel("Numerical Variables")
plt.show()
```



```
[42]: spearman_matrix = df[numerical_columns].corr(method='spearman').abs()
fig = plt.figure(figsize=(10,5))
sns.heatmap(spearman_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.title('Spearman Correlation Matrix - Numeric variables')
plt.show()
```



```
[43]: numerical_columns.remove('pdays')
spearman_matrix = df[numerical_columns].corr(method='spearman').abs()
fig = plt.figure(figsize=(10,5))
sns.heatmap(spearman_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.title('Spearman Correlation Matrix - Numeric variables')
plt.show()
```



In the spearman correlation matrix between numerical variables, we notice that there is a very strong correlation between the variables 'pdays' and 'previous', which indicates that these variables together in a model can be counterproductive in the training of a model. Therefore, we can consider deleting one of the variables.

Due to the results of the interaction analysis of each of the available variables with the target, we decided to delete the variable pdays because it maintains a very strong interaction with the target (check point biserial correlation with target).

```
[44]: # Final dataset - Next step train model
final_columns = list(set(categorical_columns)) + list(set(numerical_columns)) + \
    ['y']
print(final_columns)
final_dataset = df[final_columns].copy()
final_dataset.to_csv('./data/final_data.csv', sep='|', index=None)

['contact', 'housing', 'poutcome', 'job', 'loan', 'month', 'marital', 'default',
'education', 'age', 'previous', 'balance', 'day_of_week', 'campaign',
'duration', 'y']
```